

# Graph Neural Networks

Naeemullah Khan

[naeemullah.khan@kaust.edu.sa](mailto:naeemullah.khan@kaust.edu.sa)



جامعة الملك عبد الله  
للعلوم والتقنية

King Abdullah University of  
Science and Technology


KAUST Academy  
King Abdullah University of Science and Technology

February 1, 2023

- ▶ Complex graphs can be successfully generated via sequential generation using deep learning.
- ▶ Two of the ways that we can do this by leveraging RNNs or Reinforcement Learning.
- ▶ In GraphRNN, we sequentially predict nodes and then edge connections for these nodes.
- ▶ In GCPN (RL based method), the model predicts potential links using node embeddings.

- ▶ Recent years have seen impressive performance of deep learning models in a variety of applications.
- ▶ For example, In computer vision, deep convolutional networks have achieved human-level performance on ImageNet (image category classification).
- ▶ But are these models ready to be deployed in real world?

- ▶ Deep convolutional neural networks are vulnerable to adversarial attacks.
- ▶ For example, Imperceptible noise changes the prediction.



**Adversarial example**


Adopted from  
Goodfellow et al.  
ICLR 2015

- ▶ Adversarial examples are also reported in natural language processing [Jia Liang et al. EMNLP 2017] and audio processing [Carlini et al. 2018] domains.


- ▶ The existence of adversarial examples prevents the reliable deployment of deep learning models to the real world.
  - Adversaries may try to actively hack the deep learning models.
  - The model performance can become much worse than we expect.
- ▶ Deep learning models are often not robust.
  - In fact, it is an active area of research to make these models robust against adversarial examples.

- ▶ Common applications of GNNs involve public platforms and monetary interests.
  - Recommender systems
  - Social networks
  - Search engines
- ▶ Adversaries have the incentive to manipulate input graphs and hack GNNs' predictions.

# Graph Adversarial Learning (cont.)




↓ Prediction



“Class 2”


80.4% confidence

+




Perturbations

=



↓ Prediction

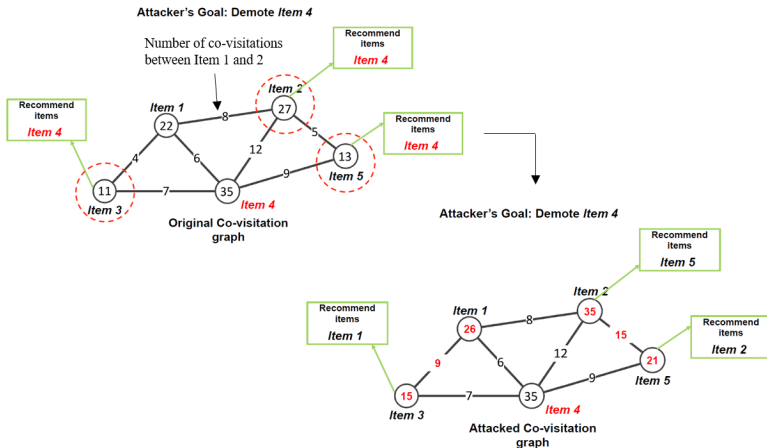


“Class 3”

92.1% confidence


- Target
- Influencer
- Class1
- Class2
- Class3
- $[:]$  Node features

# Graph Adversarial Learning (cont.)











- ▶ What are the attack possibilities in real world?
  - Target node  $t \in V$ : node whose label prediction we want to change.
  - Attacker nodes  $S \subset V$ : nodes the attacker can modify.




- ▶ **Direct Attack:** Attacker node is the target node:  $S = \{t\}$
- ▶ Modify target node feature.
  - For example, change website content.
- ▶ Add connections to target.
- ▶ Remove connections from target.
  - For example, unfollow users.



- ▶ **Indirect Attack:** The target node is not in the attacker nodes:  $t \notin S$
- ▶ Modify attacker node features.
  - For example, hijack friends of targets.
- ▶ Add connections to attackers.
  - For example, create a link, link farm.
- ▶ Remove connections from attackers.
  - For example, delete undesirable link.



- **Objective for the attacker:** Maximize (change of target node label prediction) Subject to (graph manipulation is small).
- If graph manipulation is too large, it will easily be detected. Successful attacks should change the target prediction with "unnoticeably-small" graph manipulation.



- ▶ Original graph:
  - $A$ : adjacency matrix,  $X$ : feature matrix
- ▶ Manipulated graph (aXer adding noise):
  - $A'$ : adjacency matrix,  $X'$ : feature matrix
- ▶ Assumption:  $(A, X) \approx (A', X')$ 
  - Graph manipulation is unnoticeably small i.e., preserves basic graph statistics (e.g., degree distribution) and feature statistics.
  - Graph manipulation is either direct (changing the feature/connection of target nodes) or indirect.

## Overview of the attack framework

- ▶ Original adjacency matrix  $A$ , node features  $X$ , node labels  $Y$ .
- ▶  $\theta^*$ : Model parameter learned over  $A, X, Y$ .
  - $c_v^*$ : class label of node  $v$  predicted by GCN with  $\theta^*$ .
- ▶ An attacker has access to  $A, X, Y$  and the learning algorithm.
- ▶ The attacker modifies  $(A, X)$  into  $(A', X')$ .
- ▶  $\theta^{*'}:$  Model parameter learned over  $A', X', Y$ .
  - $c_v^{*'}$ : class label of node  $v$  predicted by GCN with  $\theta^{*'}$ .
- ▶ The goal of the attacker is to make  $c_v^* \neq c_v^{*'}$ .

- ▶ Target Node:  $v \in V$ .
- ▶ GCN learned over the original graph

$$\theta^* = \arg \min_{\theta} \mathcal{L}_{train}(\theta; A, X)$$

- ▶ GCN's original prediction on the target node:

$$c_v^* = \arg \max_c f_{\theta^*}(A, X)_{v,c}$$

- ▶ GCN learned over the manipulated graph

$$\theta^{*'} = \arg \min_{\theta} \mathcal{L}_{train}(\theta; A', X')$$

- ▶ GCN's prediction on the target node:

$$c_v^{*'} = \arg \max_c f_{\theta^{*'}}(A', X')_{v,c}$$

- ▶ We want the prediction to change after the graph is manipulated:

$$c_v^* \neq c_v^{*'}$$



## Change of prediction on target node $v$ :

$$\Delta(v; A', X') =$$

$$\log f_{\theta^{*'}}(A', X')_{v, c_v^{*'}} - \log f_{\theta^{*'}}(A', X')_{v, c_v^*}$$

Predicted (log)  
probability of the  
newly-predicted  
class  $c_v^{*'}$



Want to increase  
this term




Want to decrease  
this term

- ▶ **Setting:** Semi-supervised node classification with GCN.
- ▶ **Graph:** Paper citation network (2,800 nodes, 8,000 edges).
- ▶ **Attack type:** Edge modification (addition or deletion of edges).
- ▶ **Attack budget on node  $v$ :**  $d_v + 2$  modifications ( $d_v$ : degree of node  $v$ ).
  - Intuition: It is harder to attack a node with a larger degree.
- ▶ Model is trained and attacked 5 times using different random seeds.

Predicted probabilities of a target node  $v$  over 5 re-trainings (each bar represents a single trial)


(without graph manipulation, i.e., clean graph)



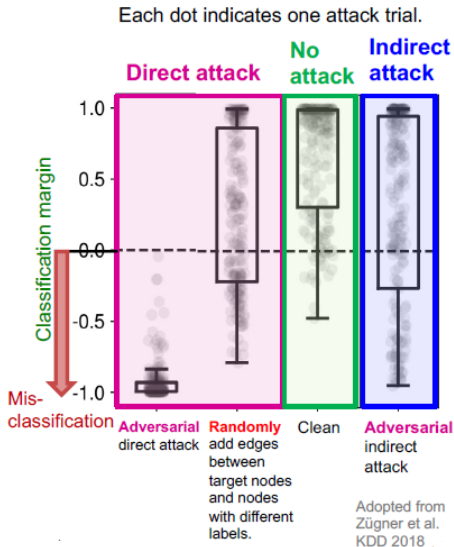
**Classification margin**  
> 0: Correct classification  
< 0: Incorrect classification

7-class classification

GCN's prediction after modifying 5 edges attached to the target node (direct adversarial attack).



- ▶ **Adversarial direct attack** is the strongest attack, significantly worsening GCN's performance (compared to no attack).
- ▶ **Random attack** is much weaker than adversarial attack.
- ▶ **Indirect attack** is more challenging than direct attack.



The defense methods can be categorized into three main categories.

- ▶ **Adversarial Training:** The model is made more resilient to adversarial attacks by training it on perturbed graphs as well as clean graphs during training.
- ▶ **Adversarial Detection:** These methods assume that the data has already been polluted and employ preprocessing methods to detect and remove/reduce the effect of attacks.
- ▶ **Robust Optimization:** These methods employ more robust loss functions to improve the robustness of models.

- ▶ Machine Learning models can be very sensitive to very perturbations and can be tricked.
- ▶ GCN's prediction performance can be significantly harmed by adversarial attacks.
- ▶ GCN is not robust to adversarial attacks but it is somewhat robust to indirect attacks and random noise.
- ▶ Several defense methods exist which are used to improve the robustness of models.

These slides have been adapted from

- ▶ Jure Leskovec, Stanford CS224W: [Machine Learning with Graphs](#)
- ▶ Stephan Gunnemann, [Graph Neural Networks: Adversarial Robustness](#)
- ▶ Jintang Li, Spectra: [Adversarial Learning on Graph Introduction](#)